

Hyperparameters tuning

Part I

Plan for the lecture

- Hyperparameter tuning in general
 - General pipeline
 - Manual and automatic tuning
 - What should we understand about hyperparameters?
- Models, libraries and hyperparameter optimization
 - Tree-based models
 - Neural networks
 - Linear models

Plan for the lecture: models

- Tree-based models
 - GBDT: XGBoost, LightGBM, CatBoost
 - RandomForest/ExtraTrees
- Neural nets
 - Pytorch, Tensorflow, Keras...
- Linear models
 - SVM, logistic regression
 - Vowpal Wabbit, FTRL
- Factorization Machines (out of scope)
 - libFM, libFFM

How do we tune hyperparameters

How do we tune hyperparameters

1. Select the most influential parameters

- a. There are tons of parameters and we can't tune all of them

How do we tune hyperparameters

1. Select the most influential parameters

- a. There are tons of parameters and we can't tune all of them

2. Understand, how exactly they influence the training

How do we tune hyperparameters

1. Select the most influential parameters

- a. There are tons of parameters and we can't tune all of them

2. Understand, how exactly they influence the training

3. Tune them!

- a. Manually (change and examine)
- b. Automatically (hyperopt, etc.)

Hyperparameter optimization software

- A lot of libraries to try:
 - *Hyperopt*
 - Scikit-optimize
 - Spearmint
 - GPyOpt
 - RoBO
 - SMAC3

Automatic hyperparameter optimization

```
def xgb_score(param):  
    # run XGBoost with parameters `param`  
  
def xgb_hyopt():  
    space = {  
        'eta' : 0.01,  
        'max_depth' : hp.quniform('max_depth', 10, 30, 1),  
        'min_child_weight' : hp.quniform('min_child_weight', 0, 100, 1),  
        'subsample' : hp.quniform('subsample', 0.1, 1.0, 0.1),  
        'gamma' : hp.quniform('gamma', 0.0, 30, 0.5),  
        'colsample_bytree' : hp.quniform('colsample_bytree', 0.1, 1.0, 0.1),  
  
        'objective': 'reg:linear',  
  
        'nthread' : 28,  
        'silent' : 1,  
        'num_round' : 2500,  
        'seed' : 2441,  
        'early_stopping_rounds': 100  
    }  
  
    best = fmin(xgb_score, space, algo=tpe.suggest, max_evals=1000)
```

Color-coding legend

1. Underfitting (bad)
2. **Good fit and generalization (good)**
3. Overfitting (bad)

Color-coding legend

- A parameter in red
 - *Increasing it impedes fitting*
 - Increase it to reduce overfitting
 - Decrease to allow model fit easier

Color-coding legend

- A parameter in red
 - *Increasing it impedes fitting*
 - Increase it to reduce overfitting
 - Decrease to allow model fit easier
- A parameter in green
 - *Increasing it leads to a better fit (overfit) on train set*
 - Increase it, if model underfits
 - Decrease if overfits

Conclusion

- Hyperparameter tuning in general
 - General pipeline
 - Manual and automatic tuning
 - What should we understand about hyperparameters?

Conclusion

- Hyperparameter tuning in general
 - General pipeline
 - Manual and automatic tuning
 - What should we understand about hyperparameters?
- Models, libraries and hyperparameter optimization
 - Tree-based models
 - Neural networks
 - Linear models